


Sapienza Training Camp 2021

Building an Image Search Engine

2 - 4 September, 2021



Roadmap



Define similarity function. Order images according to similarity to the query.

Q

Query: "person walking with a dog on the beach"



$$\text{sim}(Q, T_1) > \text{sim}(Q, T_2)$$

Highlights in Natural Language Processing (NLP)

- Google Translate
 - translate.google.com
- BERT language model used in Google search
 - [Google uses AI to boost search engine ranking efficiency, FT.com, Oct. 2019](#)
- OpenAI's GPT language model:
 - <https://openai.com/blog/better-language-models>

Recurrent neural networks

Image captioning model from:

https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/text/image_captioning.ipynb

```
class RNN_Decoder(tf.keras.Model):
    def __init__(self, embedding_dim, units, vocab_size):
        super(RNN_Decoder, self).__init__()
        self.units = units

        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(self.units,
                                       return_sequences=True,
                                       return_state=True,
                                       recurrent_initializer='glorot_uniform')


        self.fc1 = tf.keras.layers.Dense(self.units)
        self.fc2 = tf.keras.layers.Dense(vocab_size)

        self.attention = BahdanauAttention(self.units)
```

Recap: dense and convolutional layers

- Dense layer:


$$\mathbf{o} = g(W\mathbf{x} + w_0), \quad \text{where } \mathbf{x} \in \mathbb{R}^d$$



Recap: dense and convolutional layers

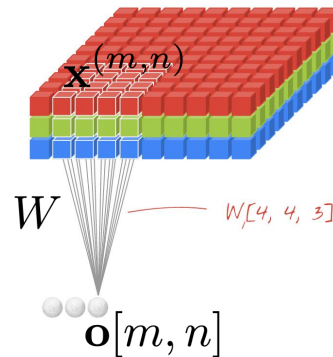
- Dense layer:

$$\mathbf{o} = g(W\mathbf{x} + w_0), \quad \text{where } \mathbf{x} \in \mathbb{R}^d$$



- Convolutional layer:


$$\mathbf{o}[m, n] = g(W\mathbf{x}^{(m, n)} + w_0), \quad \text{where } \mathbf{x}^{(m, n)} = \mathbf{x}[m:m+D, n:n+D]$$



Recap: dense and convolutional layers

- Dense layer:

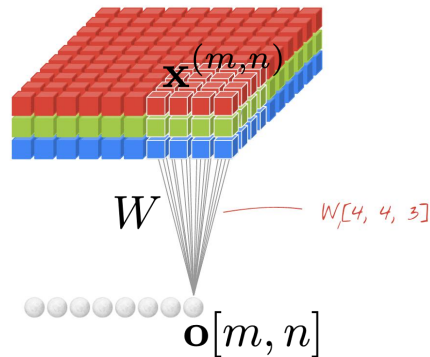
$$\mathbf{o} = g(W\mathbf{x} + w_0), \quad \text{where } \mathbf{x} \in \mathbb{R}^d$$



- Convolutional layer:


$$\mathbf{o}[m, n] = g(W\mathbf{x}^{(m, n)} + w_0), \quad \text{where } \mathbf{x}^{(m, n)} = \mathbf{x}[m:m+D, n:n+D]$$

- **weight sharing:** same weights used for all local windows $\mathbf{x}^{(m, n)}$
- convolutional layer supports variable size input



Recurrent layer

- Recurrent layer
 - **weight sharing** across time steps
 - recurrent layer supports sequences of variable size



$$\mathbf{h}^{(t)} = g(W\mathbf{h}^{(t-1)} + U\mathbf{x}^{(t)} + w_0)$$

$$\text{where } \mathbf{h}^{(t)} \in \mathbb{R}^K$$

Recurrent layer configurations




Image description (our
application!)

Recurrent layer configurations





Image description (our application!)



Text classification

Image from
["https://www.deeplearningbook.org/slides/10_rnn.pdf"](https://www.deeplearningbook.org/slides/10_rnn.pdf)

Recurrent layer configurations






Image description (our application!)



Text classification



Machine translation

Image from
https://www.deeplearningbook.org/slides/10_rnn.pdf

Image captioning model




Image captioning model

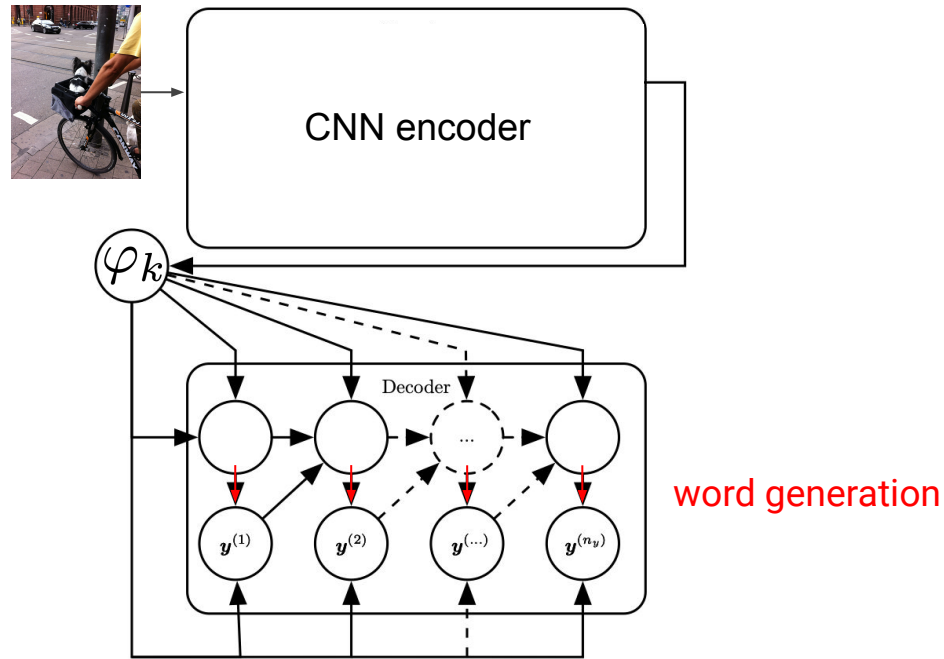


Image captioning model





Image captioning model



```
class RNN_Decoder(tf.keras.Model):  
    def __init__(self, embedding_dim, units, vocab_size):  
        super(RNN_Decoder, self).__init__()  
        self.units = units  
  
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)  
        self.gru = tf.keras.layers.GRU(self.units,  
                                       return_sequences=True,  
                                       return_state=True,  
                                       recurrent_initializer='glorot_uniform')  
  
        self.fc1 = tf.keras.layers.Dense(self.units)  
        self.fc2 = tf.keras.layers.Dense(vocab_size)  
  
        self.attention = BahdanauAttention(self.units)
```

feed generated word
back to RNN: use
embedding layer

Model training




Image from
["https://www.deeplearningbook.org/slides/10_rnn.pdf"](https://www.deeplearningbook.org/slides/10_rnn.pdf)

Model training





Image from
["https://www.deeplearningbook.org/slides/10_rnn.pdf"](https://www.deeplearningbook.org/slides/10_rnn.pdf)

Model training



Total loss:

$$\mathcal{L} = \sum_t \mathcal{L}^{(t)}$$
$$= - \sum_t \log p_{\text{RNN}}(y^{(t)} | X^{(1:t)}; \theta)$$

(compare to the cross-entropy loss on slide 17 in the “introduction to NN” module)

Image from
https://www.deeplearningbook.org/slides/10_rnn.pdf

More complex RNN units: LSTM and GRU

```
class RNN_Decoder(tf.keras.Model):
    def __init__(self, embedding_dim, units, vocab_size):
        super(RNN_Decoder, self).__init__()
        self.units = units

        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(self.units,
                                       return_sequences=True,
                                       return_state=True,
                                       recurrent_initializer='glorot_uniform')

        self.fc1 = tf.keras.layers.Dense(self.units)
        self.fc2 = tf.keras.layers.Dense(vocab_size)

        self.attention = BahdanauAttention(self.units)
```

Take a quiz!